



*A Beginner-Friendly Explanation for Aspiring
Data Scientists*

When “RAG Will Fail” in Your AI-ML-DL Project?



Sanjay N Kumar

Data scientist | AI ML Engineer | Statistician | Analytics Consultant

What Is RAG? 📖



Imagine your brain 🧠 needs to answer a question, but you forgot!

So, you quickly check a **book** 📖 and then reply.

That's what RAG does!

It **retrieves** info from a knowledge base (like Google search) and then **generates** an answer.

Simple Formula



👉 **RAG = Retrieval + Generation**

If either step fails, RAG will give a wrong answer ❌


When RAG Fails (Overview)



RAG can fail when:

1. The **retrieved data is wrong**  
2. The **retrieved data is incomplete** 
3. The **question is unclear** 
4. The **model can't connect facts properly** 
5. The **database is outdated** 

Case 1: Wrong Data Retrieved 🔍



If you ask: “Who invented the bulb?” 💡
and RAG retrieves info about *Edison’s company profits*,
the model may say something unrelated!

Real-Life Example:

You search “Lion diet” 🦁 but get results for
“Lionel Messi diet” ⚽
That’s a retrieval failure!

Mathematically:

Let R = *correct retrieval probability*.

If $R = 0.7$, then 30% of answers may fail due to bad data.

Case 2: Incomplete Information



If the database has only **half the facts**,
AI will give **half-correct answers**.

Example:

Ask: “What causes rain?” 

RAG retrieves only “clouds hold water.”

It misses “water vapor condenses and falls.”

 Result = incomplete answer.

Math View:

If dataset has only 60% coverage → expected
accuracy drops to $0.6 \times 100 = 60\%$ 

Case 3: Poor Question Understanding



If you ask, “Why do planes fly higher than birds?” 🐦 ✈️

and RAG misunderstands it as “difference between birds and planes,”
it retrieves wrong info.


Real Example:

When kids ask vague questions like “Why is sky?” 😊💧

Even humans can’t answer clearly!

So AI also struggles.

Case 4: Model Fails to Connect Facts



Sometimes AI retrieves all pieces but can't join them logically.

Example:


To answer “Why does ice float on water?”



It must connect two facts:

- Ice is less dense than water.
 - Lower density makes it float.
- If AI misses the link, answer fails.

Case 4: Model Fails to Connect Facts



Math Angle:

Let's say 2 facts = A and B.

If both retrieved (probability 0.9) but
connection fails (0.7),

then total success = $0.9 \times 0.7 = 63\%$

Case 5: Outdated Knowledge 🕒



AI may use an **old snapshot** of data.

Example:

You ask, “Who is India’s current president?”



If data is from 2020, it may give the wrong name.

Real-Life Analogy:

Like reading an old newspaper 📰 — facts have changed!

Case 6: Domain Mismatch



If RAG is trained on **medical data** 
and you ask about **cricket**, 

AI may “hallucinate” or make up facts.

That’s like asking a chef  about rocket science

Case 7: Poor Chunking or Embedding



If the data chunks are too large or small,
AI retrieves irrelevant text.

Example:

If your notes are in one big paragraph 📄,
finding one sentence is tough.

If split too tiny, context is lost.

Math Connection:

If each chunk has 10% relevant info →
retrieval accuracy $\approx 0.1 \times 100 = \mathbf{10\%}$

Case 8: Low Vector Similarity

RAG uses **vectors** (math representations) to find similar meanings.

If similarity between question vector (Q) and data vector (D) is low, retrieval fails.

Formula:

$$\text{Similarity} = \cos(\theta)$$

If $\theta = 0^\circ \rightarrow$ perfect match (1)

If $\theta = 90^\circ \rightarrow$ no match (0)

So, lower similarity $\rightarrow$ poorer answers.

Example:

If you ask “capital of USA” and data says

“Washington D.C. is a city,”

the similarity may still be low if not phrased properly.

Case 9: Memory Limit or Context Overflow



LLMs can remember only a certain amount of words (context window).

If $\text{data} > \text{limit}$, old info gets forgotten.

Example:

Like reading a 100-page book  but remembering only the last 20 pages.

Older parts vanish!

Case 10: Storage or Indexing Errors




If files are not saved or indexed properly, RAG retrieves nothing or wrong data.

Real-Life Example:

You save homework in “Downloads” but search in “Documents.”

Of course, you won't find it! 

How to Prevent RAG Failure



- ✓ Keep knowledge base updated
- ✓ Clean and organize data
- ✓ Ask clear questions
- ✓ Use good chunking (balanced size)
- ✓ Monitor retrieval accuracy with math checks


Simple Summary



RAG fails when:

- It can't **find** right info 
- It can't **understand** question 
- It can't **combine** facts properly 
- It uses **old or irrelevant** data 

Final Thought



Even the smartest AI needs **clean data + clear questions** to give **smart answers** 🗣️🤖

Just like humans — if we read the wrong book, we'll also give the wrong answer! 📖😅

Decode the Secrets of RAG Failures! 🧠

Turn confusion into clarity.

Let clean data and smart retrieval guide your AI to think, reason, and respond with precision! ⚙️ ✨

🔍 **Train your models not just to answer — but to understand.**

Reach out, and let's build intelligence that truly learns!



Sanjay N Kumar

Data scientist | AI ML Engineer | Statistician | Analytics Consultant



<https://www.linkedin.com/in/sanjaytheanalyst360/>



sanjaytheanalyst360@gmail.com